

Yue Lin



PhD Student, School of Data Science
The Chinese University of Hong Kong, Shenzhen

Blog: yuelin301.github.io
Email: linyue3h1@gmail.com
Mobile: +86-187-5817-2219

EDUCATION

- **The Chinese University of Hong Kong, Shenzhen** Shenzhen, China
PhD Student - School of Data Science 2024.8 - Present
Research Interests: Multi-Agent Reinforcement Learning, Information Design.
Fortunately advised by Prof. Baoxiang Wang and Prof. Hongyuan Zha.
Other close collaborators include Prof. Wenhao Li and Prof. Pascal Poupart.
- **Tianjin Polytechnic University** Tianjin, China
Bachelor of Engineering - Computer Science and Technology 2018.9 - 2022.6
 - **School of Computer Science and Technology** 2019.9 - 2022.6
GPA: 3.89/4 (92.2/100) Rank: 1/127
 - **School of Mechanical Engineering** 2018.9 - 2019.6
GPA: 3.90/4 (92.0/100) Rank: 1/60

EXPERIENCE

- **Microsoft** China
Research Intern @ Microsoft Research Asia (MSRA) 2026.8 - Present
Mentor: Lu Wang [Google Scholar].
- **Tencent** Shenzhen, China
Research Intern - Lightspeed Studios 2025.12 - 2026.6
Mentor: Yuan Liu.
- **The Chinese University of Hong Kong, Shenzhen** Shenzhen, China
Research Assistant - School of Data Science 2022.2 - 2024.8
Fortunately advised by Prof. Baoxiang Wang and Prof. Hongyuan Zha.

PUBLICATIONS & PREPRINTS

- Information Design in Multi-Agent Reinforcement Learning.
Yue Lin, Wenhao Li, Hongyuan Zha, Baoxiang Wang.
Neural Information Processing Systems (NeurIPS) 2023.
 - Poster. This is one of my representative works.
 - As of 2026, my advisor also regards this as one of his representative works.
 - Sources: [Paper], [Code], [Blog], [Talk].
 - TLDR: Introduces a Markov signaling game and an obedience-constrained signaling-gradient algorithm that learns what information to reveal over time so self-interested receivers are willing to follow the sender's recommendations.
- Verbalized Bayesian Persuasion.
Wenhao Li, **Yue Lin**, Yun Hua, Xiangfeng Wang, Bo Jin, Hongyuan Zha, Baoxiang Wang.
International Conference on Machine Learning (ICML) 2026.
 - Sources: [Paper], [arXiv].
 - TLDR: Introduces a verbalized Bayesian-persuasion framework in which LLMs play sender and receiver while an equilibrium solver searches for natural-language messages the receiver has reason to follow, including in multi-stage dialogues.
- Policy-Conditioned Policies for Multi-Agent Task Solving.
Yue Lin, Shuhui Zhu, Wenhao Li, Ang Li, Dan Qiao, Pascal Poupart, Hongyuan Zha, Baoxiang Wang.
arXiv preprint. 2025-12-24.
 - This work was scooped by *Evaluating LLMs in Open-Source Games*. We then planned a substantially improved version, but that follow-up direction was also scooped by Google DeepMind's *Code-Space Response Oracles: Generating Interpretable Multi-Agent Policies with Large Language Models*.
 - Sources: [Manuscript].

- TLDR: Represents each agent’s strategy as readable source code and introduces Programmatic Iterated Best Response, where an LLM rewrites one policy against another using game rewards and unit tests as feedback.
- History-Dependent Policy Gradient.
Yue Lin, Jiacheng Nie, Hongyuan Zha, Baoxiang Wang.
Draft. 2024.
 - No paper was produced; our method was scooped by Google’s paper *Multi-agent cooperation through learning-aware policy gradients (ICLR 2025)*.
 - TLDR: Derives an exact policy gradient for non-stationary multi-agent learning systems by explicitly accounting for the histories of the other agents.
- The Reciprocity Gradient.
Yue Lin, Pascal Poupart, Shuhui Zhu, Dan Qiao, Wenhao Li, Yuan Liu, Hongyuan Zha, Baoxiang Wang.
arXiv preprint. 2026-05-08.
 - This is one of my representative works.
 - Sources: [Manuscript].
 - TLDR: Introduces the reciprocity gradient, which analytically traces how actions and evaluative signals propagate through reputation chains and return as future rewards, jointly optimizing both without reward shaping.
- Talk, Judge, Cooperate: Gossip-Driven Indirect Reciprocity in Self-Interested LLM Agents.
Shuhui Zhu, **Yue Lin**, Shriya Kaistha, Wenhao Li, Baoxiang Wang, Hongyuan Zha, Gillian K. Hadfield, Pascal Poupart.
International Conference on Machine Learning (ICML) 2026.
 - Sources: [Paper], [arXiv], [Code].
 - TLDR: Introduces ALIGN, a decentralized mechanism in which self-interested language-model agents exchange open-ended, tone-graded gossip to build reputations, coordinate social norms, and isolate defectors.
- Information Bargaining: Bilateral Commitment in Bayesian Persuasion.
Yue Lin, Shuhui Zhu, William A Cunningham, Wenhao Li, Pascal Poupart, Hongyuan Zha, Baoxiang Wang.
EC 2025 Workshop: Information Economics and Large Language Models.
 - Sources: [Paper], [Code].
 - TLDR: Recasts long-term Bayesian persuasion as two-sided bargaining with commitment on both sides, separating the value of private information from the advantage of proposing first.
- Innovative Design and Simulation of a Transformable Robot with Flexibility and Versatility, RHex-T3.
Yue Lin, Yujia Tian, Yongjiang Xue, Shujun Han, Huaiyu Zhang, Wenxin Lai, Xuan Xiao.
International Conference on Robotics and Automation (ICRA) 2021.
 - Oral. Delivered a presentation at the Xi’an conference venue.
 - Sources: [Paper], [Blog].
 - TLDR: Introduces a two-degree-of-freedom transformable RHex design that switches among wheel, leg, claw, and hook modes for efficient travel, obstacle crossing, object transport, and ladder climbing.
- A snake-inspired path planning algorithm based on reinforcement learning and self-motion for hyper-redundant manipulators.
Yue Lin, Jianming Wang, Xuan Xiao, Ji Qu, Fatao Qin.
International Journal of Advanced Robotic Systems (IJARS) 2022.
 - Sources: [Paper], [Code], [Blog].
 - TLDR: Introduces Swinging Search and Crawling Control, which uses reinforcement learning over the arm’s self-motion to find a collision-free target configuration, then moves the arm toward it by crawling instead of searching every intermediate pose.

PATENTS

- Method and Apparatus for Pretraining Large Language Models Based on Federated Learning (基于联邦学习的大语言模型预训练方法和装置).
Ang Li, Baoxiang Wang, Hongyuan Zha, **Yue Lin**.
Chinese Invention Patent.
No. ZL 2026 1 0720098.0 (Grant No. CN 122242656 B), Granted Jul. 2026.
- Communication Method, Terminal Device, and Storage Medium for Multi-Agent Reinforcement Learning (多智能体强化学习通信方法、终端设备及存储介质).
Yue Lin, Wenhao Li, Hongyuan Zha, Baoxiang Wang.
Chinese Invention Patent.
No. ZL 2023 1 0397744.0 (Grant No. CN 116455754 B), Granted Sep. 2025.

PROFESSIONAL SERVICES

- NeurIPS 2025 Top Reviewer
- ICML 2026 Silver Reviewer [Award Email Screenshot]
- Independent reviewer for top-tier venues:
 - NeurIPS 2024 [6; 45615]; 2025 [5; 32912]; 2026 [4; 17738]
 - ICLR 2025 [3; 21831]; 2026 [1; 5994]
 - ICML 2025 [6; 32893]; 2026 [6; 43185]
 - AAAI 2027 [2; 12678]
 - TMLR 2025 [2; 38363]; 2026 [2; 10331]
- Volunteer:
 - ACL 2026 [1; 4155]
 - AAMAS 2024 [3, 9876], 2025 [2, 4792], 2026 [2, 2303]
 - ICML (Position) 2025 [2, 5016]
- Numbers in brackets indicate the number of manuscripts reviewed and the character count of all reviews, respectively.
Total reviews: 47.
- Served as a teaching assistant for a postgraduate course on Artificial Intelligence.